

RetailIQ

An Intelligent, Machine-Learning-Powered Retail Decision Support & Dynamic Restocking System

Project Title:	RetailIQ Intelligent Decision Support System
Authors:	Srushti Khandelwal, Simran Chawla
Domain:	Retail Analytics, Machine Learning & Supply Chain
Dataset Volume:	100,000 FMCG Transactions (24 Attributes)
Repository:	github.com/SimranChawla03/RetailIQ
Deployment:	retailiq-analytics.streamlit.app
Submission Date:	October 2026

ABSTRACT: Fast-Moving Consumer Goods (FMCG) retailers operate within razor-thin margins where poor inventory forecasting precipitates severe revenue depletion through stockouts and capital-draining overstock. RetailIQ delivers an end-to-end intelligent decision-support architecture uniting predictive multivariate regression, multi-class inventory health classification across four benchmarked models, dynamic lead-time buffer restocking logic, Apriori market-basket mining, and K-Means SKU segmentation. Engineered with vectorized NumPy execution and transparent model compression, RetailIQ delivers sub-second inference across 100,000 transaction records via a responsive Streamlit cloud interface.

Table of Contents & Executive Summary

Table of Contents

Section	Title	Page
1.0	Introduction, Background & Objectives	3
2.0	Literature Review (5 Peer-Reviewed Papers)	4
3.0	Methodology & System Mathematical Formulations	5
4.0	Data Cleaning, Preprocessing & Feature Engineering	6
5.0	Results & Empirical Discussion	7
6.0	State-of-the-Art Model Benchmark & Comparison	8
7.0	User Interface Engineering & Cloud Deployment	9
8.0	Conclusion, Future Scope & References	10

List of Figures & Tables

Identifier	Caption	Page
Table 2.1	Literature Review Comparative Synthesis Matrix	4
Figure 3.1	End-to-End System Pipeline Flowchart	5
Table 4.1	FMCG Dataset Schema & Variable Definitions	6
Table 5.1	Apriori Market Basket Association Rules	7
Table 6.1	Classification Algorithms Multi-Metric Benchmark	8
Figure 6.1	Comparative Visual Performance Across Classifiers	8
Table 7.1	Sub-Second Latency Optimization Benchmarks	9

Executive Summary & System Key Performance Indicators

The RetailIQ platform transforms transactional retail records into high-conviction inventory and revenue decisions. By replacing disjointed static heuristics with a coordinated predictive suite, the platform delivers verifiable gains across operational efficiency, capital utilization, and customer fulfillment.

0.8899 REVENUE R ² SCORE	100,000 PROCESSED TRANSACTIONS	< 0.60s PIPELINE LATENCY	83.6% MODEL COMPRESSION	148 APRIORI BUNDLE RULES
---	---	---------------------------------------	-----------------------------------	---------------------------------------

Key Architectural Innovation: RetailIQ resolves the classical deployment bottleneck of machine learning dashboards. By decoupling predictive classification from deterministic buffer replenishment and applying Gzip serialization, the production footprint was shrunk from 245 MB to 40 MB, enabling native GitHub LFS tracking and sub-second cold starts on cloud containers.

1. Introduction

1.1 Background and Industrial Context

In modern Fast-Moving Consumer Goods (FMCG) and retail operations, store managers operate in hyper-competitive markets characterized by tight gross margins (often 2% to 6%) and highly volatile consumer preferences. The profitability of physical and hybrid retail networks depends entirely on maintaining optimal stock levels. However, retail inventory management has historically suffered from the "Bullwhip Effect" and fragmented replenishment planning.

1.2 Problem Formulation

Retail businesses face two diametrically opposed financial failure modes:

- **Stockouts (Understocking):** Occur when sudden customer demand exceeds available stock. Stockouts directly cost retailers immediate turnover, cause cart abandonment, and degrade brand loyalty.
- **Overstocking:** Tying up scarce working capital in excess inventory increases warehousing carrying expenses, accelerates shrinkage, and risks catastrophic margin erosion when perishable FMCG products expire.

Legacy solutions rely on static spreadsheet reorder points (e.g., reordering when stock hits an arbitrary 50 units), completely ignoring supplier transit days, customer demographics, seasonality, and promotional basket bundling.

1.3 Specific Project Objectives

To overcome these systemic barriers, **RetailIQ** was developed with five clear objectives:

1. **Predictive Revenue Modeling:** Formulate a robust regression pipeline to anticipate line-item transaction value from volume, acquisition costs, customer demographics, and calendar cycles.
2. **Multi-Class Inventory State Classification:** Train, evaluate, and tune four distinct machine learning classifiers (KNN, Decision Tree, SVM, and Random Forest) to categorize SKUs into *Healthy Stock*, *Low Stock*, and *Overstocked*.
3. **Dynamic Lead-Time Restocking Engine:** Implement an algorithmic buffer model that calculates dynamic replenishment thresholds using empirical demand rates and vendor transit times.
4. **Data-Driven Product Intelligence:** Uncover hidden customer affinity patterns using Apriori Market Basket Analysis and group products into functional behavioral clusters using K-Means.
5. **Enterprise-Grade Production Deployment:** Engineer a zero-latency, cloud-hosted web interface featuring automated CSV normalization, automated purchase order creation, and instant export capabilities.

1.4 Scope and Boundaries

The scope encompasses 100,000 retail records from an Indian FMCG enterprise network across diverse geographical zones, formats (Hypermarket, Supermarket, Kirana), and product segments. The deliverable is a production-tested web suite capable of ingesting arbitrary retail CSVs and returning instantaneous business guidance.

2. Literature Review

To establish a solid scientific baseline, five recent peer-reviewed research papers (2021–2024) across leading journals in supply chain analytics, artificial intelligence, and software systems were examined and compared.

2.1 Critical Review of Recent Benchmark Studies

1. Punia, Nikolopoulos, Singh, & Madaan (2021) [*Int. J. Prod. Econ.*] explored a multi-stage machine learning framework for retail demand forecasting. Their research proved that tree-based ensembles (Random Forest and Gradient Boosting) reduce Mean Absolute Percentage Error (MAPE) by 18.4% over classical ARIMA. However, their paper noted significant compute latency when evaluating high-cardinality catalogs without vectorized code paths.

2. Chen, Zhang, & Liu (2022) [*IEEE Trans. Eng. Manage.*] conducted a comparative evaluation of supervised algorithms for inventory state detection. Evaluating SVM, KNN, Decision Trees, and Random Forests, they established that Random Forest delivered the most resilient macro-F1 scores under class imbalance. Conversely, Support Vector Machines suffered from extreme computational scaling constraints ($\mathcal{O}(N^3)$) on large operational datasets.

3. Al-Sharman & Bakir (2023) [*Expert Syst. Appl.*] benchmarked association rule mining algorithms in grocery e-commerce. While FP-Growth generated frequent itemset trees rapidly, the Apriori formulation provided superior operational interpretability. They demonstrated that applying a minimum Lift threshold ≥ 1.05 effectively filtered out coincidental co-purchases.

4. Kumar, Singh, & Sharma (2023) [*Comput. Ind. Eng.*] investigated SKU clustering paired with dynamic lead-time buffer stock sizing. Partitioning retail catalogs using K-Means clustering prior to computing safety stocks decreased stockout frequencies by 26% while reducing inventory carrying costs by 14%.

5. Wang, Martinez, & Gupta (2024) [*J. Syst. Archit.*] examined the memory footprint of machine learning decision support systems. They discovered that uncompressed Random Forest pickle files regularly exceed 200 MB due to redundant leaf node pointers. The authors demonstrated that serialization compression and RAM caching dropped page reload latency by 98% on cloud dashboards.


2.2 Literature Review Synthesis Matrix

Study	Focus Domain	Core Models	Key Findings	RetailIQ Integration
Punia et al. (2021)	Demand Forecasting	Random Forest, ARIMA	18.4% MAPE reduction over statistical baselines	Vectorized inference engine with pre-computed models
Chen et al. (2022)	Inventory State Class.	SVM, KNN, DT, RF	Random Forest won on imbalanced classes (F1: 0.84)	StandardScaler pipeline + Tuned Random Forest selection
Al-Sharman (2023)	Market Basket Rules	Apriori vs. FP-Growth	Lift > 1.05 guarantees high-conversion bundles	Pre-computed Apriori table cached for instant querying
Kumar et al. (2023)	Dynamic Restocking	K-Means, Buffer Sizing	26% decrease in retail stockout frequency	Dynamic lead-time formula linking daily demand buffers
Wang et al. (2024)	Cloud DSS Optimization	Model Pruning, Gzip	Sub-second response on cloud web frameworks	Gzip compression (213MB → 34.8MB) + RAM caching

3. Methodology

3.1 End-to-End System Architecture

RetailIQ follows a decoupled, sequential data science architecture. Ingested raw files flow through automated normalization, clean data splitting, simultaneous multi-model inference, dynamic reorder quantity derivation, and presentation formatting.



3.2 Mathematical Formulations

1. Revenue Regression Formulation:

Expected revenue is formulated as a linear combination of volume, pricing, customer age, loyalty, and time indices:

$$\hat{Y}_{\text{Revenue}} = \alpha_0 + \alpha_1 X_{\text{Units}} + \alpha_2 X_{\text{Cost}} + \alpha_3 X_{\text{Price}} + \alpha_4 X_{\text{Month}} + \alpha_5 X_{\text{Day}} + \alpha_6 X_{\text{Age}} + \alpha_7 X_{\text{Loyalty}}$$

2. Inventory Health Target Formulation:

Inventory health classes are derived from the empirical ratio between available stock and threshold levels ($R_{\text{stock}} = \text{ext}\{\text{Stock_On_Hand}\} / \text{ext}\{\text{Reorder_Level}\}$):

$\text{ext}\{\text{State}\} = \begin{cases} \text{Low Stock}, & \text{if } R_{\text{stock}} < 3.27 \\ \text{Healthy Stock}, & \text{if } 3.27 \leq R_{\text{stock}} \leq 8.35 \\ \text{Overstocked}, & \text{if } R_{\text{stock}} > 8.35 \end{cases}$

3. Vectorized Dynamic Buffer Restocking Logic:

Rather than using static reorder constants, RetailIQ evaluates lead-time buffer demand dynamically:

$\alpha_{D_{b,c}} = \frac{1}{N_{b,c}} \sum_{i=1}^{N_{b,c}} \text{ext}\{\text{Units}_i\}, \quad S_{\text{req}} = \alpha_{D_{b,c}} \times \text{ext}\{\text{Lead_Time_Days}\}$
 $\text{ext}\{\text{Reorder Condition}\} = \begin{cases} \text{Restock Now}, & \text{if } S_{\text{current}} \leq \text{ext}\{\text{Reorder_Level}\} \\ \text{Restock Soon}, & \text{if } S_{\text{current}} < S_{\text{req}} \\ \text{Stock Sufficient}, & \text{otherwise} \end{cases}$
 $Q_{\text{order}} = \max\left(0, \lceil S_{\text{req}} \rceil - S_{\text{current}} \right)$

4. Market Basket Association Formulation (Apriori):

Support, Confidence, and Lift are evaluated across binary basket transaction matrices:

$\text{ext}\{\text{Support}\}(A \implies B) = \frac{|\sigma(A \cup B)|}{|T|}, \quad \text{ext}\{\text{Confidence}\}(A \implies B) = \frac{|\sigma(A \cup B)|}{|\sigma(A)|}, \quad \text{ext}\{\text{Lift}\}(A \implies B) = \frac{\text{ext}\{\text{Confidence}\}(A \implies B)}{\text{ext}\{\text{Support}\}(B)}$

4. Data Cleaning and Preprocessing

4.1 Dataset Characteristics

The operational dataset encompasses **100,000 real-world retail transactions** from an Indian retail network. The schema incorporates 24 distinct attributes combining financial, operational, and customer demographics.

4.2 Dataset Attribute Dictionary

Feature	Type	Physical Meaning	Sample Values
Invoice_ID	String	Unique transaction alphanumeric key	INV10001, INV10002
Invoice_Date	Date	Timestamp of retail purchase	2024-01-15, 2024-03-22
City	Categorical	Retail store urban location	Mumbai, Delhi, Bengaluru
Store_Format	Categorical	Store physical footprint	Hypermarket, Supermarket
Category	Categorical	Department classification	Dairy, Beverages, Personal Care
Brand	Categorical	Product manufacturing brand	Amul, Britannia, Tata, HUL
Units	Integer	Physical quantity sold in transaction	1, 3, 12
Cost_Price	Float	Wholesale unit acquisition cost (₹)	45.00, 120.50
Selling_Price	Float	Retail unit price charged to buyer (₹)	60.00, 155.00
Revenue	Float	Gross line turnover (Units × Selling_Price)	180.00, 1860.00
Stock_On_Hand	Integer	Units physically present in store inventory	120, 450, 15
Reorder_Level	Integer	Baseline static reorder threshold	50, 100, 30
Lead_Time_Days	Integer	Supplier delivery transit period	3, 7, 14
Customer_Age	Integer	Customer demographic age	24, 42, 65
Loyalty_Flag	Binary	Customer loyalty membership (0 / 1)	0, 1

4.3 Automated Column Alias Normalization Engine

A primary challenge in enterprise software deployment is header inconsistency across diverse point-of-sale systems. RetailIQ features an automated alias engine in `src/preprocessing.py` containing an alias dictionary mapping over 80 variations:

- `Units` maps from: `quantity`, `qty`, `qty_sold`, `sales_qty`, `volume`.
- `Brand` maps from: `product`, `product_name`, `item`, `sku_name`.
- `Cost_Price` maps from: `unit_cost`, `purchase_price`, `buying_price`.

Headers are automatically lowercased, trimmed, and normalized via regex (`[^a-z0-9]+` → `_`) to guarantee zero-crash ingestion.

4.4 Missing Value Imputation and Feature Engineering

Rigorous integrity checks ensure data consistency: numerical values are validated for non-negativity, missing fields are imputed or computed (e.g. `Revenue = Units * Selling_Price`), and timestamps are decomposed into `Month` and `Day_of_Week` to capture seasonal sales cycles. Continuous features are scaled using `StandardScaler` for distance-sensitive algorithms.

5. Results and Discussion

5.1 Revenue Forecasting Results

The multivariate Linear Regression model demonstrated high predictive fidelity across the testing split, confirming that customer age, transaction volume, pricing, and temporal indicators reliably explain retail sales value:

- **Coefficient of Determination (R^2): 0.8899** (Explains $\approx 89\%$ of revenue variance).
- **Mean Absolute Error (MAE): ₹ 72.23** per transaction.
- **Root Mean Squared Error (RMSE): ₹ 99.09**.

5.2 Dynamic Restocking and Operational Purchase Orders

Applying the dynamic restocking engine to the 100,000-record dataset uncovered critical store imbalances:

- **Restock Now (Urgent Action):** 24,950 transactions involved products operating at or below their static reorder threshold.
- **Restock Soon (Watchlist):** Flagged items where available stock fell below the calculated lead-time demand buffer (S_{req}).
- **Purchase Order Generation:** The engine automatically aggregated replenishment quantities by Brand and Category, producing downloadable vendor purchase order CSVs.

5.3 Market Basket Association Rules (Apriori Analysis)

Apriori mining across transactional receipts discovered **148 actionable association rules** ($Lift \geq 1.05$, $Confidence \geq 30\%$). Store managers can leverage these pairings for adjacent shelf placement and promotional bundles:

Antecedent (Purchased)	Consequent (Recommended)	Support	Confidence	Lift
Dettol Hand Wash	Milk	0.051	65.38%	1.30
Olive Oil	Eggs	0.057	64.04%	1.28
Ashirvaad Whole Wheat Atta	Conditioner	0.050	60.24%	1.25
MTR Ready-to-Eat Poha	Rice	0.050	60.24%	1.24
Fresh Apples	Rice	0.062	60.19%	1.24

5.4 Product & Customer Segmentation (K-Means)

K-Means clustering with $K=9$ achieved the highest Silhouette Score (**0.1898**), outperforming Agglomerative Hierarchical clustering (**0.1630**). The model identified clear operational segments: *High-Volume Daily Staples*, *High-Margin Gourmet Goods*, and *Volatile Extended-Lead Items*.

6. State-of-the-Art Models Comparison

6.1 Multi-Model Benchmarking Framework

To identify the optimal classifier for retail inventory health, four state-of-the-art architectures were benchmarked under identical train/test splits across both **Baseline** and **Tuned** stages:

6.2 Empirical Performance Metric Comparison

Model Architecture	Hyperparameter Stage	Accuracy	Precision	Recall	F1-Score	Inference Latency
K-Nearest Neighbors (KNN)	Baseline	0.4532	0.3782	0.4533	0.3802	High (Distance scans)
Decision Tree (CART)	Baseline	0.3752	0.3782	0.3752	0.3766	Very Low (< 0.05s)
Support Vector Machine (SVM)	Baseline	0.5009	0.2509	0.5009	0.3343	Extreme (> 12.0s)
Random Forest Classifier	Baseline	0.4696	0.3766	0.4696	0.3700	Moderate (~ 0.8s)
K-Nearest Neighbors (KNN)	Tuned (k=15, Manhattan)	0.4252	0.3799	0.4252	0.3898	High
Decision Tree	Tuned (Max_Depth=8)	0.3752	0.3782	0.3752	0.3766	Very Low
Support Vector Machine	Tuned (C=1.0, RBF)	0.5009	0.2509	0.5009	0.3343	Extreme
Random Forest (Champion)	Tuned (N=50, Max_Depth)	0.4652	0.3822	0.4652	0.3778	Low (0.21s in RetailIQ)

6.3 Visual Metric Distribution Across Tuned Models



6.4 Critical Model Trade-Off Analysis

- **Support Vector Machine:** While exhibiting an apparent accuracy of 50.09%, SVM collapsed into assigning solely the majority class, yielding a disastrous precision of 0.2509. Its $\mathcal{O}(N^3)$ computational scaling makes it impractical for operational deployment.
- **Random Forest (Champion):** Achieved the highest multi-class Precision (38.22%) and balanced classification across all three inventory states. Ensemble tree voting effectively eliminated single-tree variance.

7. User Interface Development and Deployment

7.1 Consolidated Dashboard Architecture

To eliminate interface redundancy, the RetailIQ dashboard was restructured into four dedicated operational hubs:

- **Overview Hub:** High-level business KPIs, monthly sales trends, and department revenue breakdowns.
- **Inventory & Restock Hub (Unified):** Integrates stock health distribution, immediate reorders, lead-time watchlist, catalog status, and 1-click vendor purchase order CSV export.
- **Sales & Forecast Hub:** Linear Regression actual vs. predicted revenue tracking and brand turnover predictions.
- **Product Intelligence Hub:** Interactive Apriori cross-sell combo selectors and K-Means customer/product clustering.

7.2 Performance Engineering & Latency Benchmarks

Executing an ML pipeline over 100,000 records initially resulted in a 9.5-second lag. Through systematic engineering optimizations, end-to-end execution dropped to **under 0.6 seconds**:

Pipeline Component	Before Optimization	After Optimization	Speedup Factor
Model Ingestion	2.50 – 4.00 s (Disk re-read)	0.00 s (@st.cache_resource)	Instantaneous
Inventory Classification	5.79 s (4 models)	0.21 s (Champion-only)	27.5x faster
Restocking Engine	0.76 s (df.apply loops)	0.08 s (Vectorized NumPy)	9.5x faster
Clustering & Sales	0.31 s	0.28 s	1.1x faster
Total Pipeline Latency	~ 9.50 s	< 0.60 s	16x faster

7.3 Model Compression and Cloud Deployment

To comply with GitHub's strict 100 MB file limit and enable zero-lag cloud container deployment, models were compressed using Gzip Level 3 serialization:

- random_forest.pkl : Reduced from **213.19 MB** → **34.86 MB** (83.6% reduction).
- knn.pkl : Reduced from **12.90 MB** → **5.22 MB**.
- decision_tree.pkl : Reduced from **5.03 MB** → **0.75 MB**.
- **Total Repository Model Footprint:** Reduced from ~ **245 MB** down to ~ **40 MB**, tracked via Git LFS.
- **Cloud Deployment:** Successfully hosted on Streamlit Community Cloud (retailiq-analytics.streamlit.app).

8. Conclusion & Future Work

8.1 Summary of Contributions

RetailIQ bridges the gap between predictive retail modeling and operational supply chain decision-making. The project successfully delivers:

- An end-to-end revenue forecasting model explaining 89% of transaction variance ($R^2 = 0.8899$).
- A multi-class inventory health classifier evaluated across four benchmark algorithms, selecting tuned Random Forest as champion.
- A dynamic lead-time buffer replenishment engine generating itemized purchase orders.
- Apriori market basket mining discovering 148 product bundling combinations.
- A 16x faster execution pipeline handling 100,000 records in sub-second latency (< 0.60 s) deployed in the cloud.

8.2 Future Research Directions

1. **Deep Learning Sequence Modeling:** Implement Temporal Fusion Transformers (TFT) for multi-horizon seasonal forecasting.
2. **Automated ERP Integration:** Link generated purchase orders directly to supplier EDI protocols.
3. **Dynamic Markdown Pricing:** Integrate price elasticity models to automate discount markdowns on overstocked perishable SKUs.

9. References

1. Punia, S., Nikolopoulos, K., Singh, S. P., & Madaan, J. (2021). Predictive analytics in retail: Multi-stage machine learning framework for demand forecasting and stock allocation. *International Journal of Production Economics*, 234, 108035.
2. Chen, Y., Zhang, H., & Liu, M. (2022). Comparative analysis of supervised learning algorithms for multi-class inventory state detection. *IEEE Transactions on Engineering Management*, 69(4), 1120–1132.
3. Al-Sharman, A., & Bakir, M. (2023). Accelerating association rule mining in modern e-commerce and supermarket retail. *Expert Systems with Applications*, 215, 119342.
4. Kumar, R., Singh, R. K., & Sharma, A. (2023). Integrated SKU segmentation and dynamic buffer restocking in FMCG supply chains. *Computers & Industrial Engineering*, 178, 108990.
5. Wang, L., Martinez, C., & Gupta, P. (2024). Edge-cloud architectures and model pruning for real-time retail decision support systems. *Journal of Systems Architecture*, 148, 103078.
6. Agrawal, R., Imieliński, T., & Swami, A. (1993). Mining association rules between sets of items in large databases. *ACM SIGMOD Record*, 22(2), 207–216.
7. Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
8. Raschka, S. (2018). MLxtend: Providing machine learning and data science utilities and extensions to Python's scientific computing stack. *Journal of Open Source Software*, 3(24), 634.